

# Jaweria Fayyaz

✉ jawerifayyaz474@gmail.com | ☎ 03320475160 | 🔗 linkedin.com/in/jaweria-fayyaz | 🌐 github.com/jawerifayyaz

## SUMMARY

Data Engineer with experience delivering production-grade batch ETL/ELT pipelines and cloud-native analytics on Azure Databricks for P&G. Skilled in PySpark, Delta Lake, and Medallion Architecture, with hands-on experience in data migration, CDM harmonization, and retail analytics across multi-country enterprise environments.

## EDUCATION

### BS Computer Science

2021 – 2025

Govt. College University, Lahore — CGPA: 3.70

## PROFESSIONAL EXPERIENCE

### Associate Data Engineer

Aug 2025 – Mar 2026

Symufolk, Lahore

- Designed and maintained scalable batch and streaming data pipelines in Azure Databricks using PySpark, processing large-scale datasets for analytical workloads.
- Developed PySpark transformation jobs in Databricks, improving data quality through automated validation rules across multiple datasets.
- Optimized ETL workflows reducing average pipeline runtime by 25%, improving daily processing efficiency.
- Contributed to migration of legacy data from SharePoint into ADLS Gen2, building Bronze–Silver–Gold Medallion pipelines with Common Data Model harmonization.

## PROJECTS

### SHARP HVS Analytics Pipeline — P&G Retail POS

- Built a PySpark pipeline in Azure Databricks to process retail POS data from 6 markets (UAE, Saudi Arabia, Kuwait, Qatar, Oman, Thailand), normalizing weekly and monthly retailer data using a 4-4-5 fiscal calendar into a unified weekly grain.
- Engineered a statistical HVS detection algorithm using median, standard deviation, and percentile thresholds with two-level config-driven parameter lookup, classifying bulk sales per product-store-retailer and isolating clean organic consumer demand.
- Delivered a Gold Layer output table aggregating organic consumer sales monthly across LuLu, Carrefour, and Panda retailers, consumed by P&G analytics and Power BI dashboards.

### Azure Data Lakehouse & Analytics Pipeline

- Built an end-to-end data lakehouse on Azure using ADLS Gen2 and Delta Lake implementing Medallion Architecture (Bronze, Silver, Gold).
- Developed ETL/ELT pipelines using Azure Data Factory and Databricks (PySpark) for scalable data processing and transformation.
- Designed data models and optimized query performance in Azure Synapse Analytics using indexing and distribution strategies.

### Anxiety Detection (Final Year Project) — [Link]

- Built an ML/NLP web app to detect anxiety presence, type, and severity from Urdu text; created a custom dataset of 12,000+ samples with tokenization, stopword removal, and TF-IDF preprocessing.
- Trained SVM, LightGBM, XGBoost, and XLM-RoBERTa; finalized DistilBERT achieving 91.31% accuracy; deployed via Flask (backend) and React (frontend).

## CERTIFICATIONS

- Introduction to Data Engineering on Azure — Microsoft

Sep 2025

## SKILLS

**Languages:** Python, SQL, C++, Java

**Data Engineering:** ETL/ELT Pipelines, PySpark, Azure Data Factory, ADLS Gen2, Databricks, Delta Lake, dbt, Apache Airflow, Medallion Architecture, Data Modeling, Query Optimization

**Cloud Platforms:** Microsoft Azure (ADF, ADLS Gen2, Synapse Analytics, Unity Catalog), Google Cloud Platform (BigQuery, Cloud Storage), Snowflake

**Data Science:** Machine Learning, Natural Language Processing, Pandas, NumPy, Scikit-learn, NLTK, Matplotlib, Seaborn

**Databases:** MS SQL Server, MySQL, PostgreSQL

**Tools:** Git, GitHub, VS Code, Jupyter, Google Colab, Azure DevOps